

2020.04.13

Data handling

Real-time systems

Real-time tasks and scheduling

Juan Antonio de la Puente juan.de.la.puente@upm.es

Real-Time Requirements

- Some actions have to be executed within a specified time interval
 - ▶ too early or too late is wrong even though the functional behaviour is correct
- Different levels of criticality
 - ▶ **hard** RT requirements: have to be met always
 - ▶ **soft** RT requirements: may fail occasionally
 - ▶ **firm** RT requirements: if not met the result is useless

Task model

T = period
D = deadline
C = processor time
R = response time

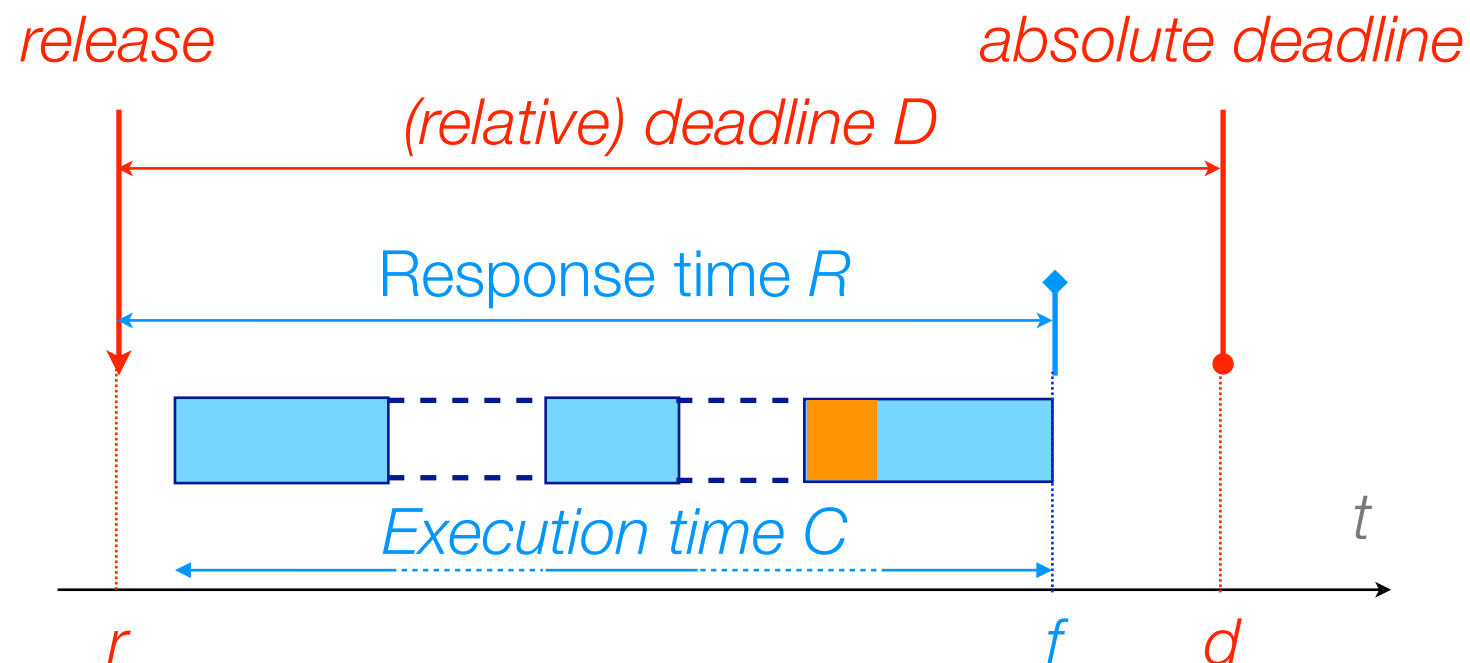
- An RTS is composed of a number of *tasks*
 - ▶ each task executes a sequence of *jobs*, which are released repeatedly according to an *activation pattern*
 - periodic, sporadic, aperiodic, random, etc.
 - ▶ each job has a deadline D relative to its release time
 - ▶ each consumes a certain amount C of processor time
- The time when the job completes, relative to its release is the response time R of the job
- The main real-time requirement is that for all jobs of every task $R \leq D$
 - ▶ $C \leq R$ is implied

$$R \text{ (RESPONSE)} \leq D \text{ (DEADLINE)}$$

$$C \text{ (processor time)} \leq R \text{ (RESPONSE)}$$

Feasible interval


- A *job* has to be executed between its release time and its deadline
- It uses processor time, possibly split into several execution intervals



Periodic tasks

T = period
 D = deadline
 C = processor time
 R = response time

- The jobs of a periodic task are released at regular intervals, with a period T
- A periodic task is characterised by (T, C, D)
 - ▶ T : period
 - ▶ C : computation time
 - ▶ D : deadline



Sporadic tasks

- The jobs of a periodic task are released whenever some event E occurs
 - ▶ a minimum separation T is usually imposed
- A sporadic task is characterised by (T, C, D)


T = period

D = deadline

C = processor time

R = response time

E = event




Aperiodic tasks

- Aperiodic tasks are event-driven, but have no real-time requirements
 - ▶ no minimal separation between event occurrences
 - ▶ no deadline
 - but there is often the requirement that jobs must complete as soon as possible

System dynamics and real-time

- Real-time requirements are derived from the system dynamic characteristics
- Example: attitude control
 - ▶ a change in a reaction wheel momentum makes the attitude angles change
 - ▶ but the change is not instantaneous



Sampling rate

sampling rate $\geq 2 \cdot \text{bandwidth}$
 $\Delta t < \Delta R / 10$

- Analog variables are sampled when reading them into a computer
- The sampling rate must be at least twice the bandwidth
 - ▶ Nyquist-Shannon sampling theorem
 - but faster sampling rates are most common
- Heuristic rules are often used
 - ▶ a common rule is $\Delta t < \Delta R / 10$
- Control functions are applied on a periodic basis
 - ▶ the period is derived from the sampling rate used in the control algorithm

Deadlines

- Deadlines are also derived from system dynamics
- A function must be completed in time for the system state to evolve as specified
- “Too late” is often not admissible
 - ▶ the system may enter an undesirable state or become uncontrollable
 - ▶ telecommand / telemetry messages may be lost
 - ▶ a ground station may lose track of a satellite

Sample requirements

- Nominal mode attitude control
- Keep nominal attitude
 - ▶ adjust roll, pitch and yaw angles by acting on reaction wheels
 - every 200 ms
 - ▶ calibrate gyro sensors with IR sensor
 - every 1s for 15 min, twice a day
 - ▶ gyro sensors deliver data every 100 ms
- Other tasks
 - ▶ dump momentum from reaction wheels
 - whenever necessary
 - ▶ send TM data when requested by TC
 - not more often than 60 ms
 - ▶ respond to TC
 - within 190 ms

Modes of operation

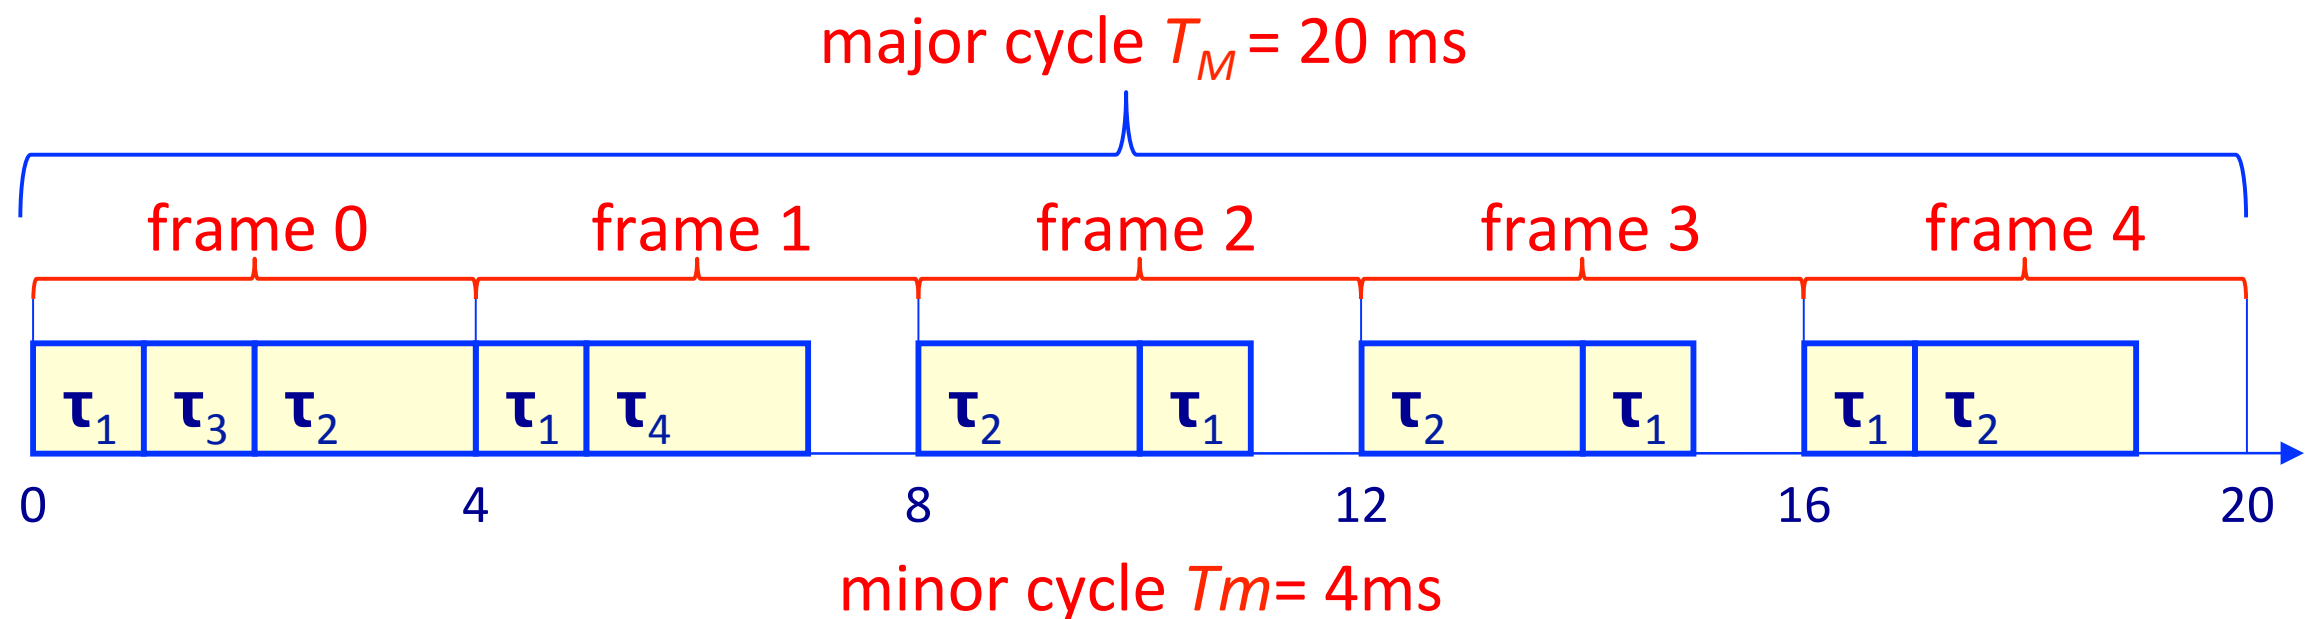
- Many real-time systems have different modes of operation
 - ▶ for example, ground, launch, orbit acquisition, nominal
- The set of tasks that are executed, as well as their temporal requirements, may vary from one mode to another
- Mode changes are triggered by specific events
 - ▶ e.g. getting to nominal orbit
- They must be carried out in such a way that the temporal integrity of the system is not compromised
 - ▶ no deadline misses, no inadmissible release delays

Time-driven implementation

- Let us consider a simple cyclic task model
 - ▶ static set of tasks
 - ▶ all tasks are periodic
- A static execution scheme can be devised for this simple model
 - ▶ the overall system behaviour is periodic
 - all actions are repeated with a hyperperiod $H = \text{mcm } T_i$
 - time is divided into (major) cycles with a duration equal to a hyperperiod
 - major cycles are divided into frames (minor cycles) of equal duration
 - individual jobs are executed within frames
 - ▶ the above scheme can be implemented with a simple program called a cyclic executive

Example

Task	T	C	D	f0	f1	f2	f3	f4
τ_1	4	1	4	x	x	x	x	x
τ_2	5	2	5	x		x	x	x
τ_3	20	1	20	x				
τ_4	20	2	20		x			



Cyclic executive

- Each frame is synchronised with a timer interrupt

```
procedure Cyclic_Executive is
  type Frame is mod 5;
  f : Frame := 0;
begin
  loop
    Wait_Timer_Interrupt; -- low-level handler
    case f is
      when 0 => T1; T3; T2;
      when 1 => T1; T4;
      when 2 => T2; T1;
      when 3 => T2; T1;
      when 4 => T1; T2;
    end case;
    f := f + 1;
  end loop;
end Cyclic_Executive;
```

Assessment

- Advantages
 - ▶ simple, robust, easy to verify
 - ▶ deterministic time behaviour
- Disadvantages
 - ▶ static schedule may be complex to build (NP-hard)
 - easier if periods are harmonic
 - long periods are difficult to handle
 - long jobs may have to be split into shorter segments
 - ▶ too rigid, difficult to maintain
 - small changes may force the whole plan to be re-computed
 - ▶ sporadic tasks are not easily incorporated into a cyclic plan
 - polling servers are commonly used

Event-driven implementation


- Tasks are implemented as threads managed by a multiprogramming kernel
 - ▶ e.g. Ada tasks
- Threads await an activation event, then execute their job and suspend until the next activation event
 - ▶ may be a clock-driven synchronous event (periodic tasks) or an asynchronous event (aperiodic/sporadic tasks)
 - ▶ see examples in chapter on clocks and timers
- Deadline guarantees are provided by scheduling and static analysis
 - ▶ next chapter

Example: periodic task pattern

```
task type Periodic (Period : Time_Span);  
  
task body Periodic is  
    Next_Time : Time := Clock;  
begin  
    loop  
        delay until Next_Time;  -- activation event  
        ... -- periodic job  
        Next_Time := Next_Time + Period;  
    end loop;  
end Periodic;
```

Processor scheduling

- If the number of concurrent tasks is greater than the number of processors, a decision must be made about which task(s) to execute at a given time
- A scheduling algorithm is used to this purpose by the OS/RTS



Some common scheduling methods

- **Round-robin scheduling**
 - ▶ each task runs for the duration of a time-slice and then gives way to the next task in a circular queue
 - ▶ common in general-purpose OS
 - good for fairness
 - ▶ poor performance in RTS
- **Priority scheduling**
 - ▶ each task has a priority
 - a measure of its importance or urgency
 - ▶ the highest-priority ready task is dispatched for execution when a processor is available
 - ▶ can provide predictable behaviour in real-time systems

Priority scheduling


- Static priorities
 - ▶ task priorities do not change over time
 - ▶ better predictability
- Dynamic priorities
 - ▶ task priorities may change
 - ▶ better processor utilisation
- Pre-emptive vs non-pre-emptive scheduling
 - pre-emptive scheduling: if a task with a higher priority than the currently running task becomes ready, the latter is pre-empted from the processor in order to run the highest-priority task
 - non-preemptive scheduling: the running task is allowed to run until it suspends itself, even if a higher priority task has become ready

Fixed-priority pre-emptive scheduling

- FPPS is a common scheduling method for real-time systems
 - ▶ task priorities are fixed
 - except perhaps on mode changes and other discrete events
 - ▶ the highest priority active task is always running
- Problem: how to assign priorities to tasks
 - ▶ Deadline-monotonic scheduling (DMS) is optimal
 - tasks with shorter deadlines have higher priorities
 - if $D = T$ for all tasks, it is the same as rate-monotonic scheduling (RMS)

Example

Task	T	C	D	P
τ_1	10	2	10	2
τ_2	20	4	8	3
τ_3	30	5	20	1



Task scheduling in Ada

- Tasks have priorities

```
subtype Any_Priority is Integer range ...;  
subtype Priority is Any_Priority  
  range Any_Priority'First .. value;  
subtype Interrupt_Priority is Any_Priority  
  range Priority'Last + 1 .. Any_Priority'Last;
```

- The **basic priority** of tasks is defined by

```
task type T  
  with Priority => P  
is ...
```

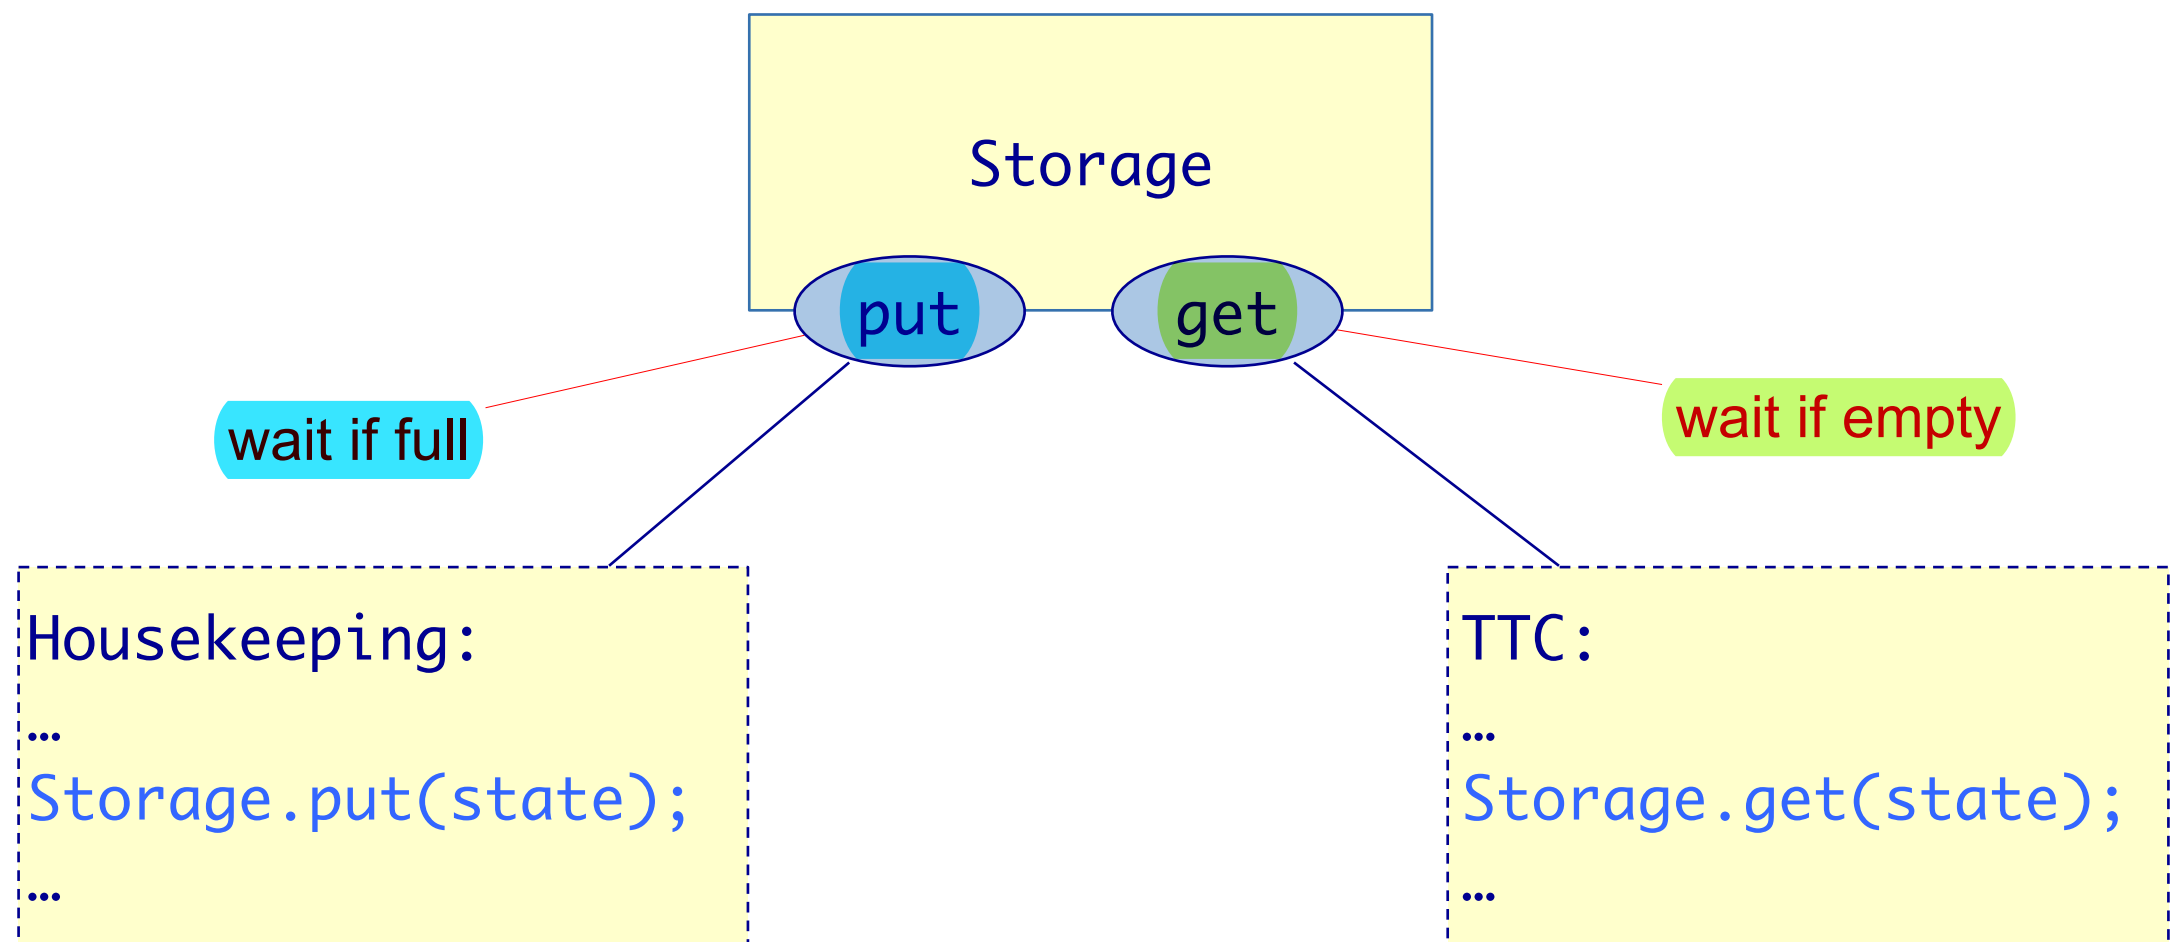
- It can be a task attribute

```
task type T (P : System.Priority)  
  with Priority => P  
is ...
```

```
Housekeeping : T(25);
```


Shared data

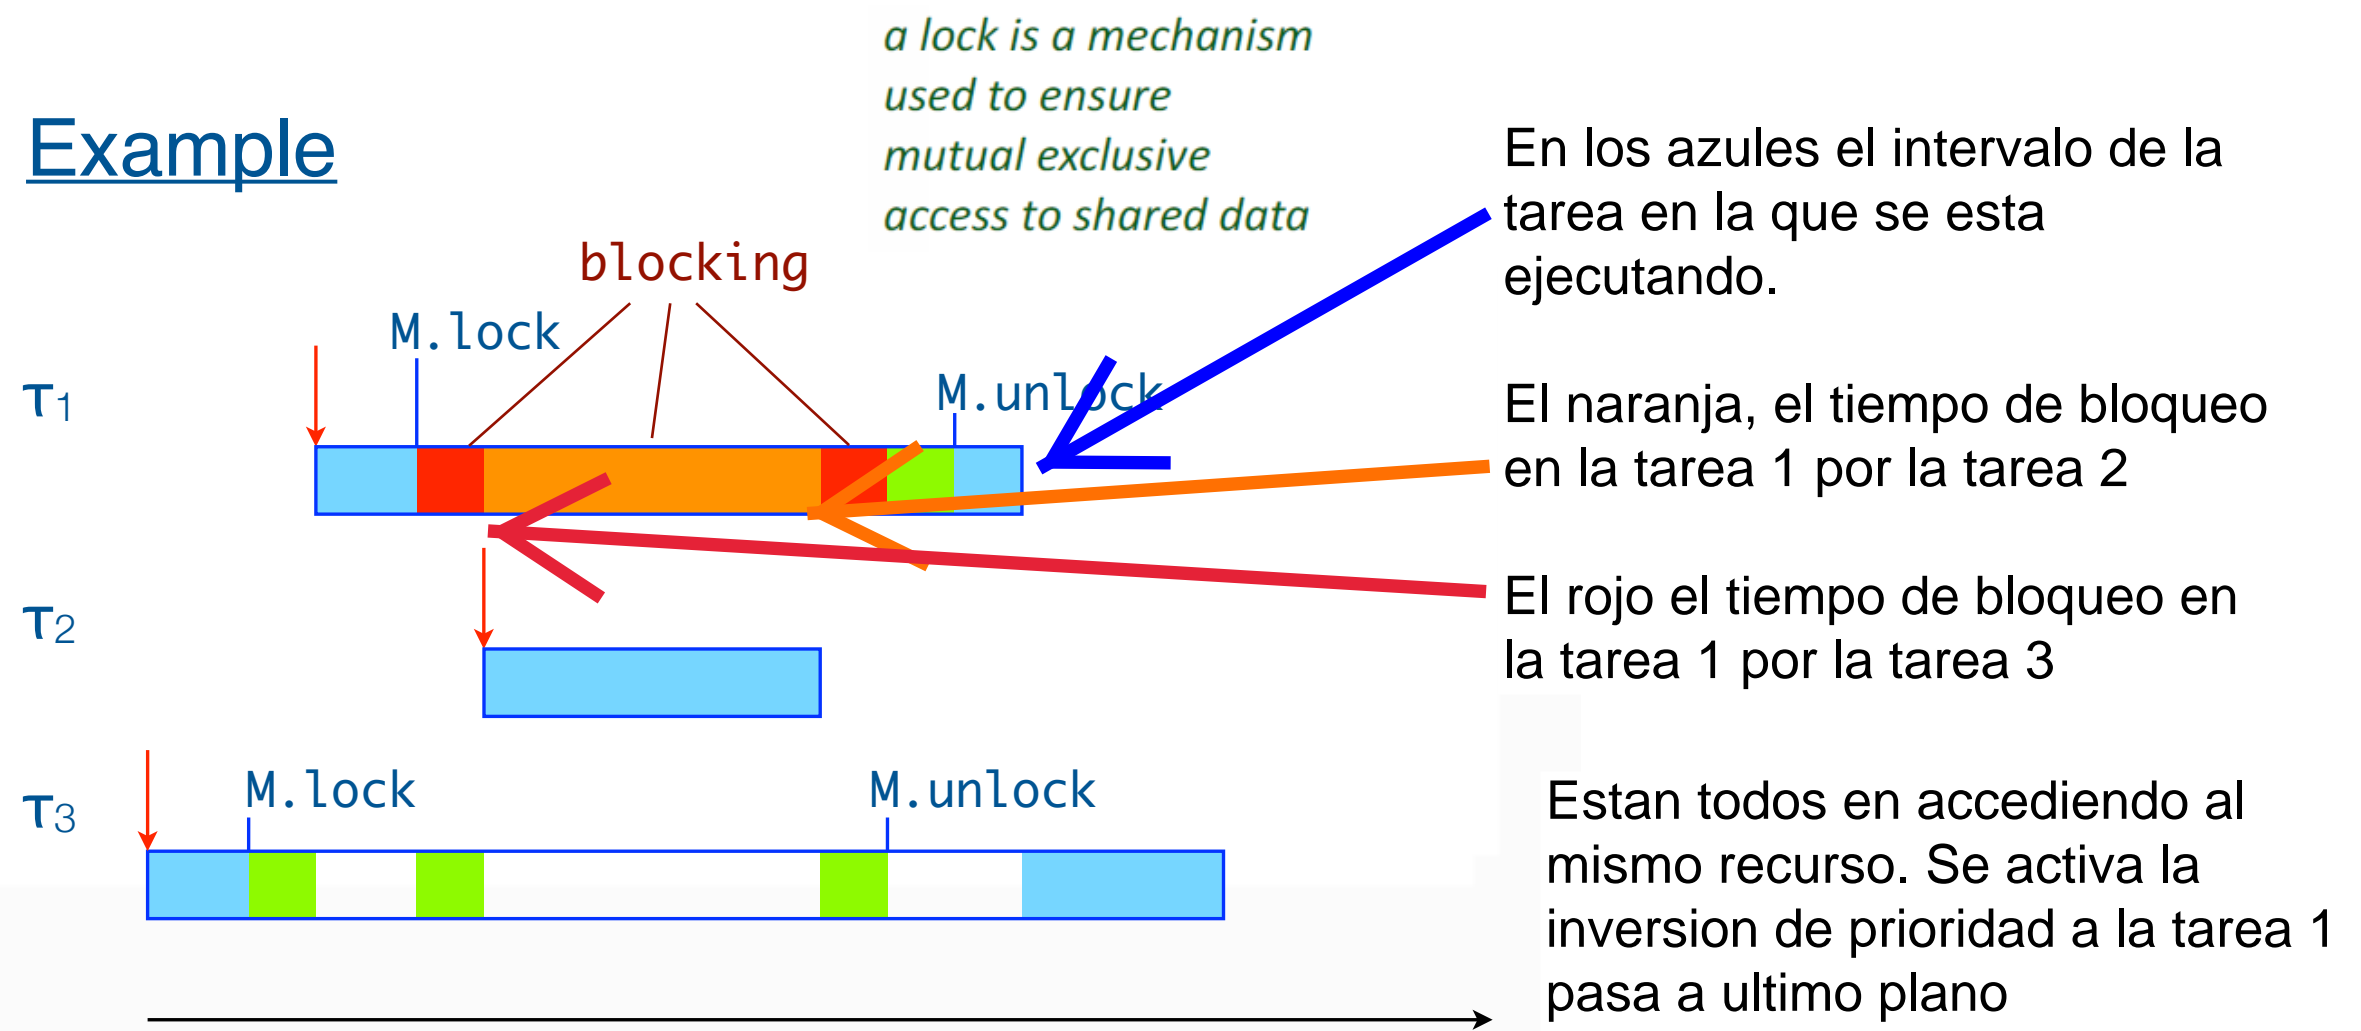
- When tasks access a shared object in mutual exclusion priority inversion may arise



Priority inversion

- When two tasks access shared data in mutual exclusion, priority inversion may arise
 - ▶ a high-priority task is delayed by a low-priority task


▶ Example





Access protocols

- Priority inversion cannot be totally eliminated
- But it can be bounded by using an appropriate access protocol for shared objects
- The immediate ceiling priority protocol (ICPP) is often used
 - ▶ each shared resource has a ceiling priority
 - the highest of the tasks that use it
 - ▶ a task trying to get exclusive access to a shared object immediately inherits its ceiling priority
 - ▶ the blocking a task may incur is bounded by the longest operation on the shared object
 - ▶ also known as *ceiling locking* or *highest locker protocol*


Example



without ceiling locking



with ceiling locking



Ceiling priority

- The ceiling priority of protected objects can be specified with an annotation

```
protected type Manager (P : System.Priority)
  with Priority => P
is
  ...
private
  ...
end Manager;

System_Manager : Manager(30);
```

Scheduling methods in Ada

- Scheduling is defined for the whole program

- ▶ in a configuration file

- ▶ FPPS is specified by

```
pragma Task_Dispatching_Policy (FIFO_Within_Priorities);
```

- ▶ ICPP for shared objects is specified by

```
pragma Locking_Policy (Ceiling_Locking);
```

- Other scheduling methods are also available

- ▶ dynamic priorities, earliest-deadline first, round-robin, etc.

